

Training on

Splunk User & Basic Administration

Splunk Overview

What is Splunk?

Splunk is a big data tool.

It is software that helps in

- Collecting machine data (logs)

- Indexing

- Searching

Helps in identifying patterns, detecting anomalies.

Helps in making better business/investment and operational decisions.

splunk® > enterprise

Bigdata Tool

splunk® > cloud™

Bigdata Tool on Cloud



SIEM

splunk® >
phantom

SOAR

Splunk Architecture

Splunk Basic Architecture



Users



Search Head



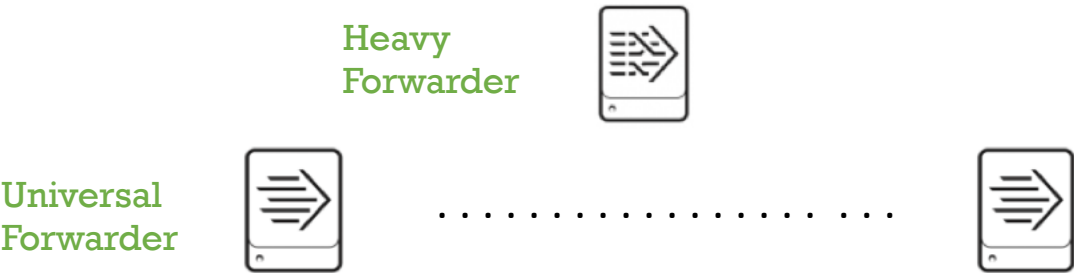
1 - 3

Indexer



1 - 5

Forwarder



1 - 10

20 - 300

Log Sources



Universal forwarder

Cannot Filter or Mask data

Doesn't have UI

Installed on the log sources itself

Can collect performance data

Heavy forwarder

Can Filter and Mask data

Has UI/Full Splunk instance

Standalone server

Cannot collect performance data

Splunk Deployment Options

Option 1



Universal Forwarder



Heavy Forwarder



Indexer



Search Head

Option 2



Universal Forwarder



Indexer
Heavy Forwarder



Search Head



Indexer
Search Head

Option 3

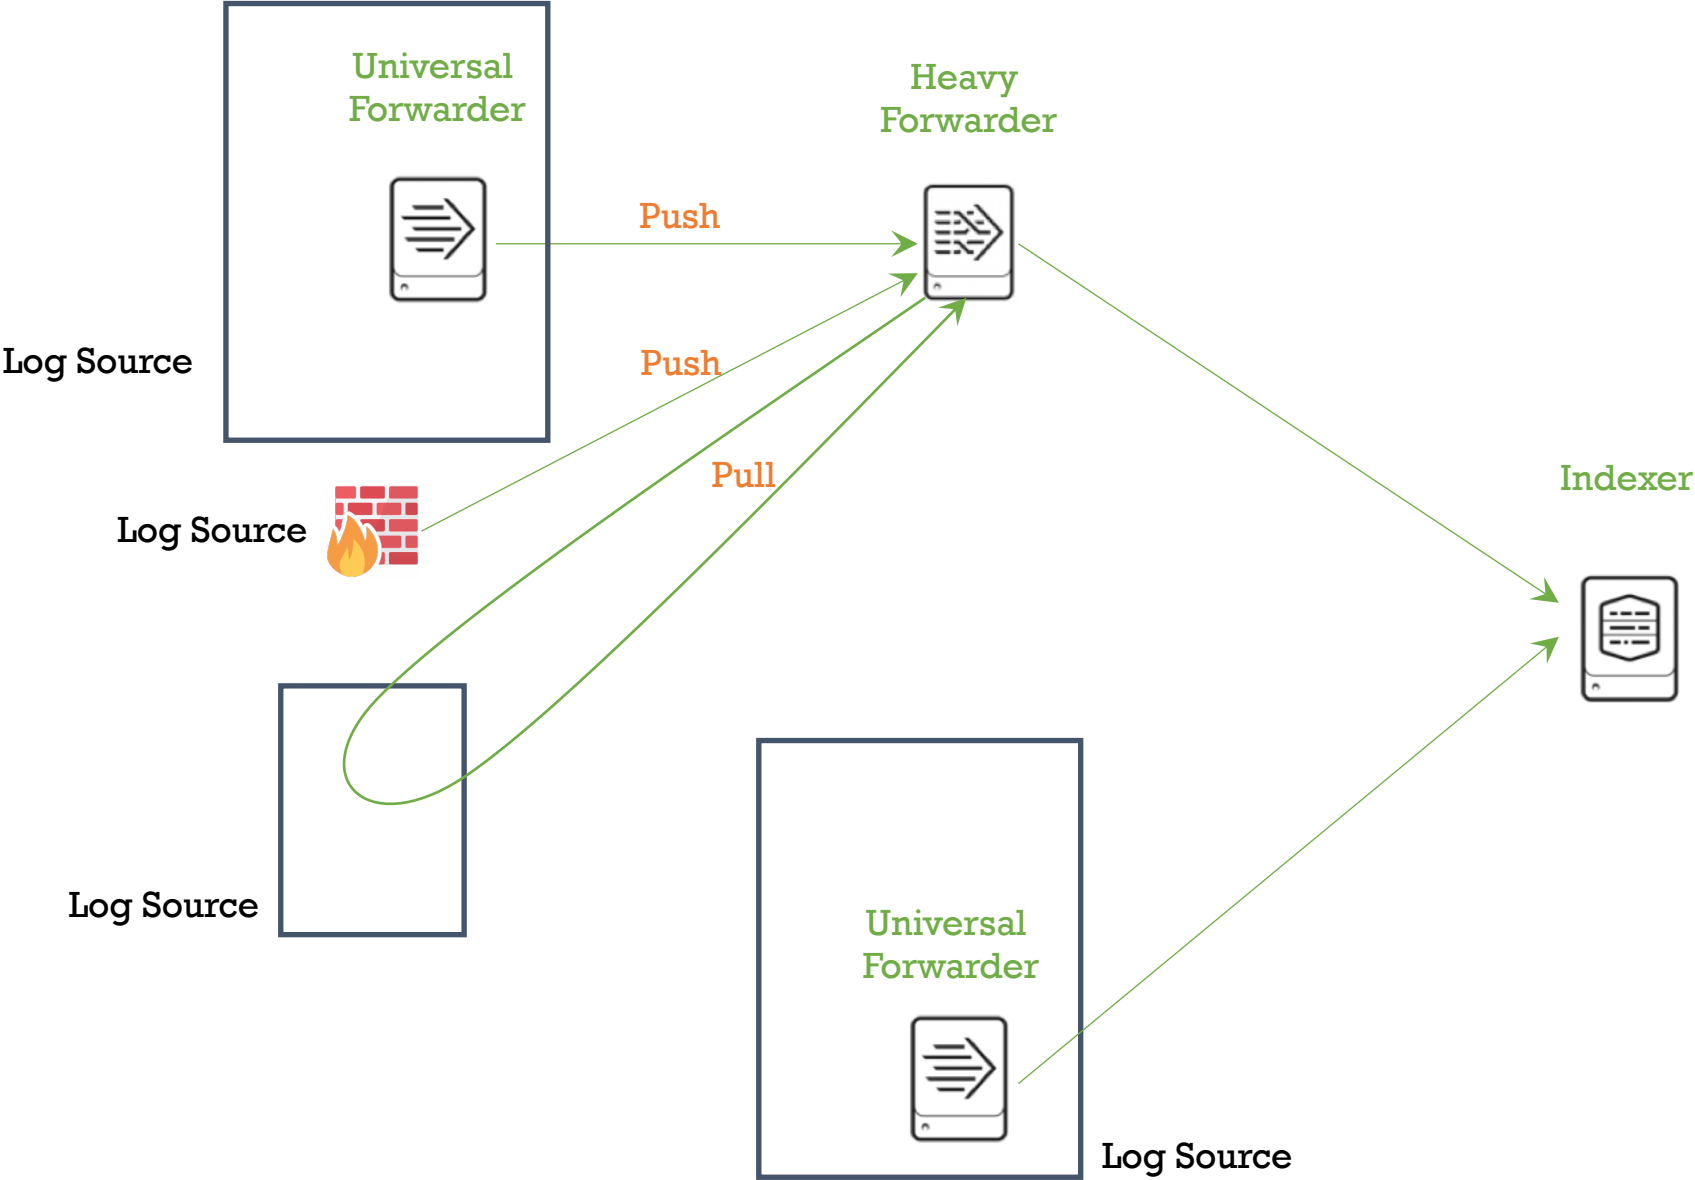


Universal Forwarder



Indexer
Heavy Forwarder
Search Head

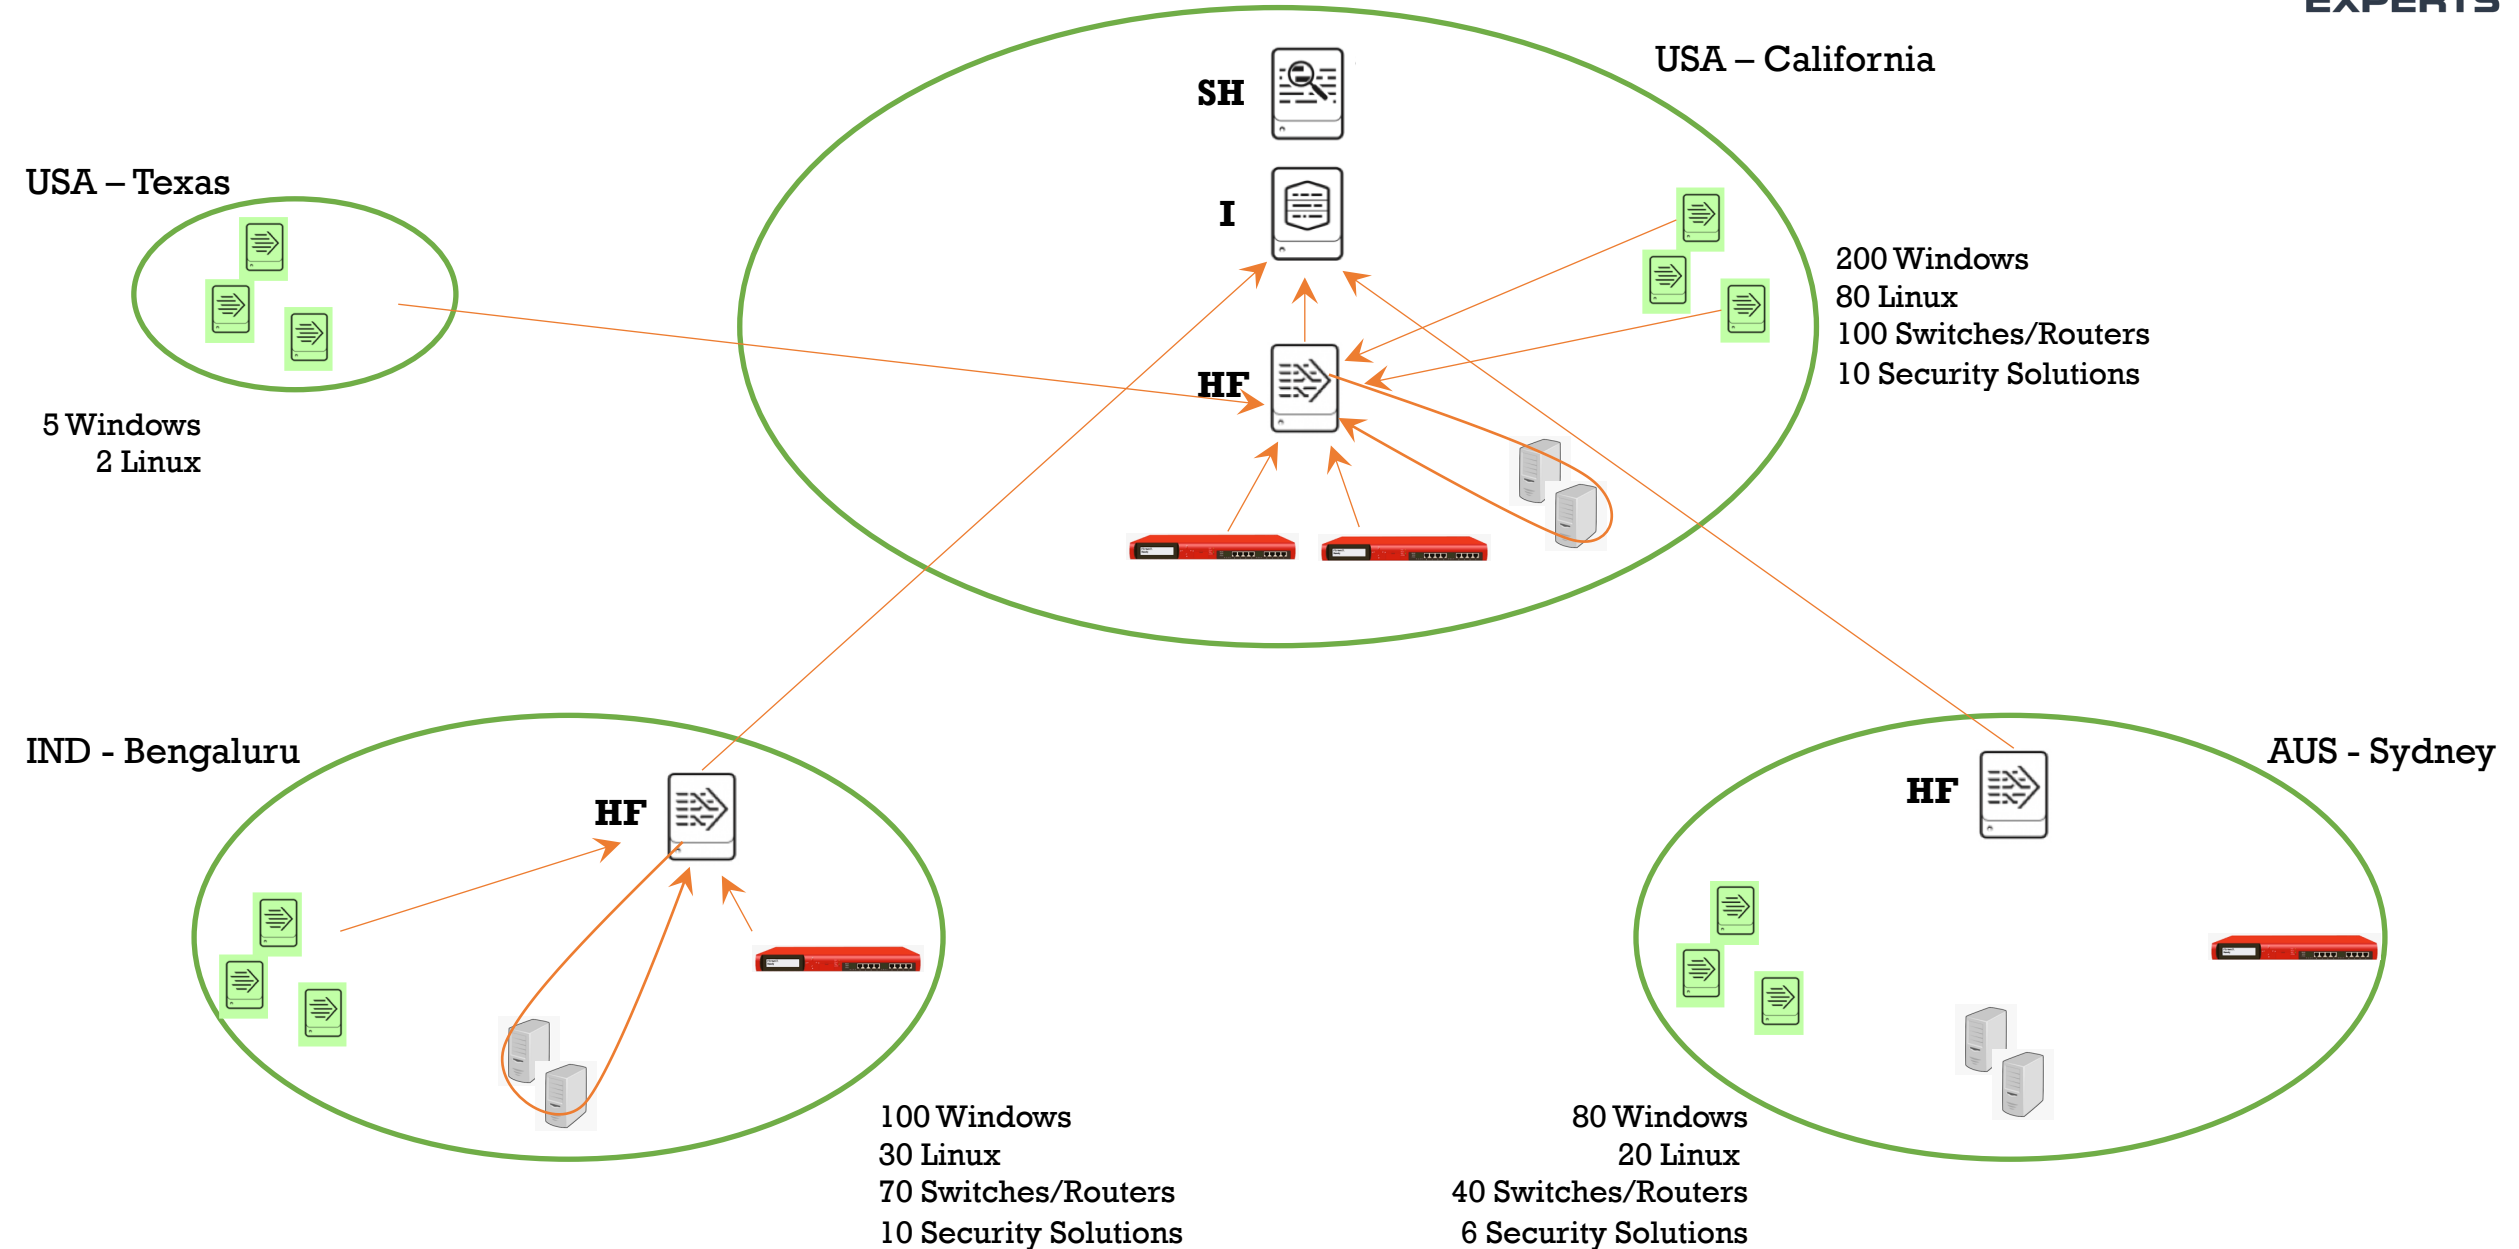
Data Flow



Activity

Example Splunk Deployment

Example Splunk Deployment



Splunk Licensing

Data indexed per day 10 GB/day 50 GB/day 100 GB/day 1 TB/day

EPS – Events Per Second

EPS x 400 bytes x 60 sec x 60 min x 24 hours

2000 x 400 bytes x 60 sec x 60 min x 24 hours = 6912000000 bytes = 69.12 GB

For Windows Event Logs we assume 1 event = 1000 bytes

Storage requirement will be calculated as **Data/Day x Retention-Time**

Splunk Installation

Splunk Installation Packages

Splunk Enterprise



Heavy Forwarder



Indexer




Search Head

Splunk Universal Forwarder




Universal Forwarder


Splunk Enterprise - Supported Platforms



Windows 10
Windows Server 2016, 2019



Linux
(CentOS/RedHat/Ubuntu/Debian)



Mac OS

Splunk UF - Supported Platforms



Windows



Linux



Solaris



Mac OS



FreeBSD



AIX

Splunk Enterprise – Port Requirements



Management Port	: 8089	(Splunk Component to Component)
Indexing Port	: 9997	(UF/HF to Indexer for log forwarding)
Web Access Port	: 8000	(User Computer to Search Head)
Syslog Port	: 514	(Log Source to HF)

Splunk Enterprise - Installation



Install Amazon Linux

Update Amazon Linux

Install wget

Download Splunk Enterprise package

Extract the package

Start Splunk service

Open port 8000 on AWS instance

Start Splunk service on boot

Install CentOS

```
#sudo yum update
```

```
#sudo yum install wget
```

```
#sudo wget <splunk-package-link>
```

```
#sudo tar xvzf splunk-package
```

```
splunk/bin/# sudo ./splunk start --accept-license*
```

Security → Security Group → Edit Inbound rules

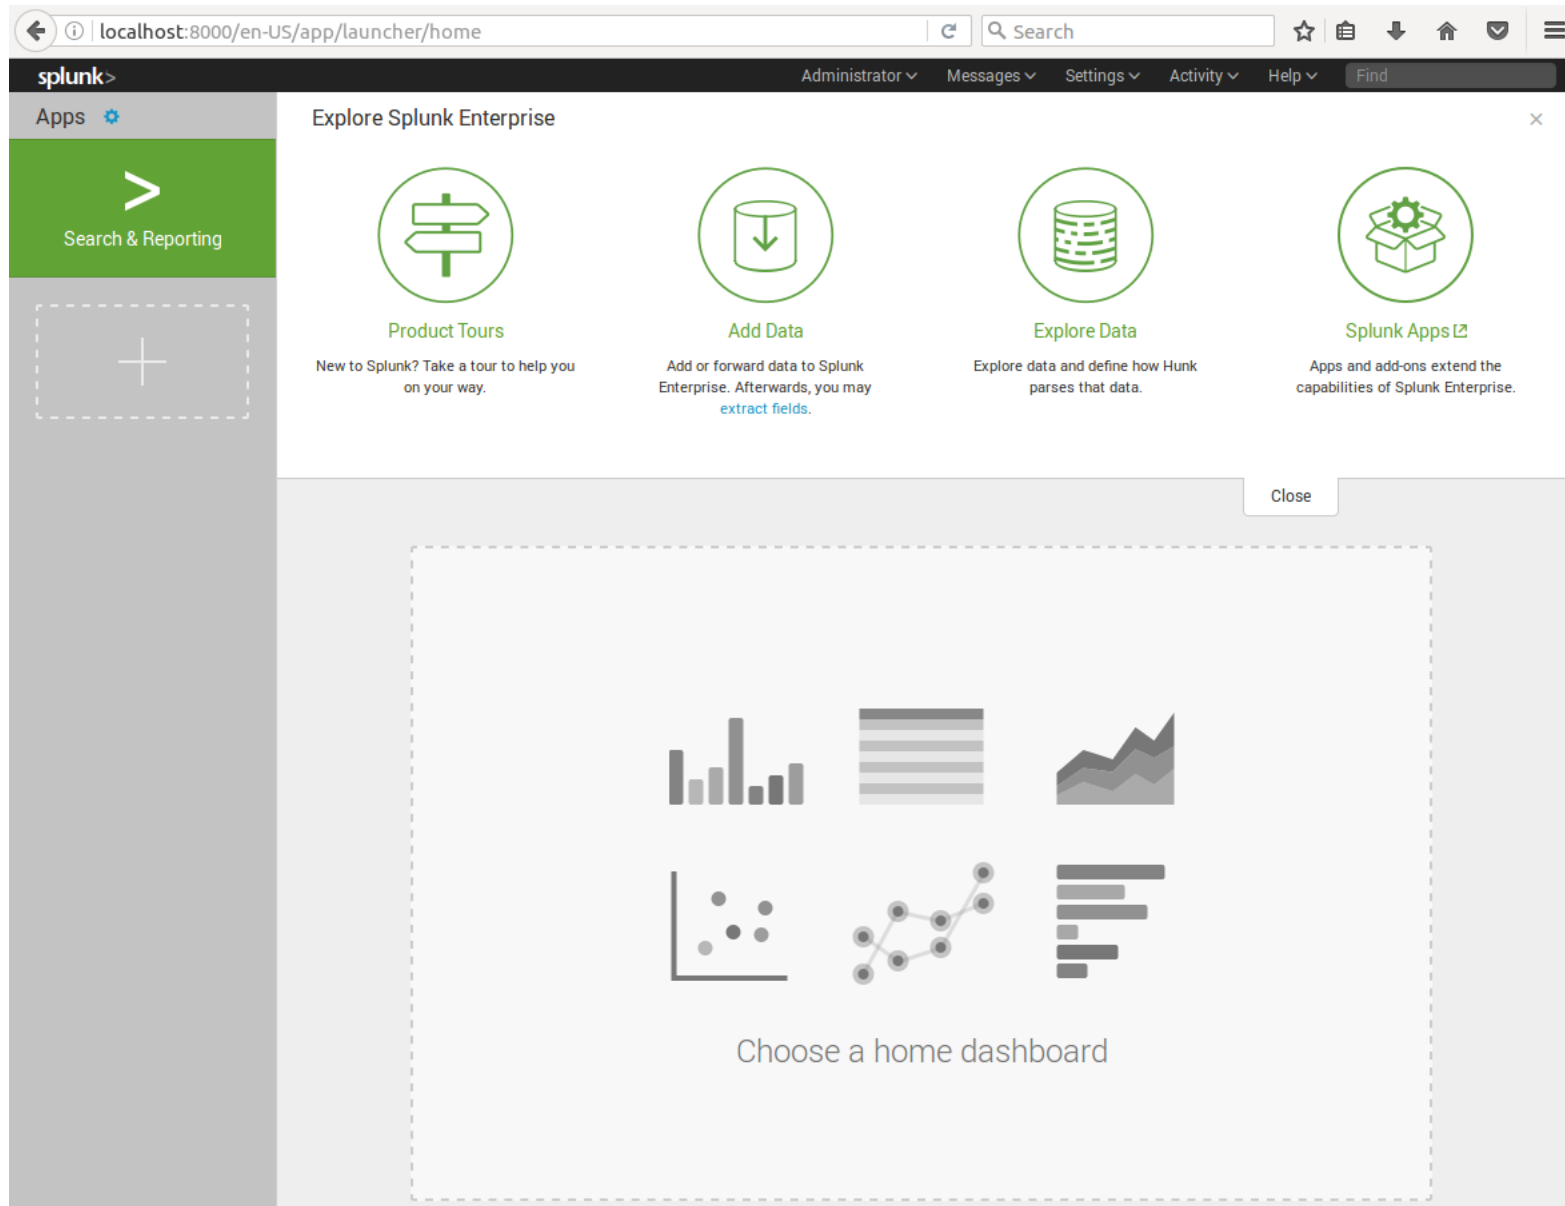
```
splunk/bin/#sudo ./splunk enable boot-start
```

* Need to set username and password for Splunk Web Access

- Installation is a one-time activity in an org.
- You cannot clear Splunk Admin or Architect interviews with knowledge acquired from our training.
- You might join a company where Splunk is not used.
- We aim to become SOC analysts or Security analysts (No SIEM Admin)
- You might encounter a few interviews regarding the Splunk Admin role – THEY ARE NOTE FOR YOU
- If you are not following Splunk – **IT IS FINE**
- Basic commands are enough (**80-20 rule**)

Accessing the Web Interface

Splunk can be accessed by the URL: http://splunk_server_ip:8000



Demo

Splunk UF – Installation (Collecting logs from Windows)

Splunk UF – Installation (Collecting logs from Ubuntu)



Install Ubuntu

Install CentOS

Update Ubuntu

`#apt update`

Install wget

`#apt install wget`

Download Splunk Universal Forwarder package

`#wget <splunk-package-link>`

Extract the package

`#tar xvzf splunk-package`

Start Splunk service

`splunk/bin/#./splunk start --accept-license*`

Start Splunk service on boot

`splunk/bin/#./splunk enable boot-start`

Demo

Splunk UF – Installation (Collecting logs from Ubuntu)

Splunk UF – Configuration Files



What to collect?

/opt/splunkforwarder/etc/system/local/**inputs.conf**

Where to send?

/opt/splunkforwarder/etc/system/local/**outputs.conf**

inputs.conf

```
[monitor:///var/log/auth.log]
```

```
index = linux
```

outputs.conf

```
[tcpout:send-to-indexer]
```

```
server = <indexer_ip>:9997
```


Search and Reporting App

The screenshot shows the Splunk Search and Reporting App interface. At the top is a dark navigation bar with the Splunk logo, the app name 'Search & Reporting', and user/setting links. Below this is a secondary navigation bar with tabs for Search, Metrics, Datasets, Reports, Alerts, Dashboards, and a 'Search & Reporting' button. The main content area is titled 'Search' and contains a search input field, a time range selector, and two informational panels: 'How to Search' and 'What to Search'. A 'Search History' section is at the bottom. Eight red circles with numbers 1 through 8 are overlaid on the interface, with arrows pointing to specific elements.

Navigation Bar: splunk>enterprise App: Search & Reporting Administrator Messages Settings Activity Help Find

Secondary Navigation Bar: Search Metrics Datasets Reports Alerts Dashboards **3** Search & Reporting

Search Section:

- 1**: Alerts tab
- 2**: Search & Reporting button
- 3**: Dashboards tab
- 4**: Search input field (enter search here...)
- 5**: Time range selector (Last 24 hours)

How to Search **6**

If you are not familiar with the search features, or want to learn more, see one of the following resources.

[Documentation](#) [Tutorial](#)

What to Search **7**

109,864 Events	9 days ago	a day ago
INDEXED	EARLIEST EVENT	LATEST EVENT

[Data Summary](#)

8: Search History

Search UI Options

Number	Element	Description
1	Applications menu	Switch between Splunk applications that you have installed. The current application, Search & Reporting app, is listed. This menu is on the Splunk bar.
2	Splunk bar	Edit your Splunk configuration, view system-level messages, and get help on using the product.
3	Apps bar	Navigate between the different views in the application you are in. For the Search & Reporting app the views are: Search, Metrics, Datasets, Reports, Alerts, and Dashboards.
4	Search bar	Specify your search criteria.
5	Time range picker	Specify the time period for the search, such as the last 30 minutes or yesterday. The default is Last 24 hours .
6	How to search	Contains links to the <i>Search Manual</i> and the <i>Search Tutorial</i> .
7	What to search	Shows a summary of the data that is uploaded on to this Splunk instance and that you are authorized to view.
8	Search history	View a list of the searches that you have run. The search history appears after you run your first search.

Search Results

Number of search matches

Event Distribution Over Time

Events (384,383)

Patterns

Statistics

Visualization

Format Timeline

Zoom Out

Zoom to Selection

Deselect

1 day per column

List

Format

50 Per Page

Prev

1

2

3

4

5

6

7

8

...

Next

< Hide Fields

All Fields

SELECTED FIELDS

a host 6

a source 17

a sourcetype 14

INTERESTING FIELDS

a attack 100+

a class 22

date_hour 14

date_mday 7

date_minute 60

a date_month 3

date_second 60

i	Time	Event
>	03/08/2020 12:04:38.000	Aug 3 12:04:38 ip-172-31-18-143 systemd[1]: Stopping Create final runtime dir for shutdown pivot root... host = ip-172-31-18-143 source = /var/log/syslog sourcetype = syslog
>	03/08/2020 12:04:38.000	Aug 3 12:04:38 ip-172-31-18-143 systemd[1]: Stopping D-Bus System Message Bus... host = ip-172-31-18-143 source = /var/log/syslog sourcetype = syslog
>	03/08/2020 12:04:38.000	Aug 3 12:04:38 ip-172-31-18-143 systemd[1]: Stopping Regular background program processing daemon... host = ip-172-31-18-143 source = /var/log/syslog sourcetype = syslog
>	03/08/2020 12:04:38.000	Aug 3 12:04:38 ip-172-31-18-143 systemd[1]: Stopped target Network is Online. host = ip-172-31-18-143 source = /var/log/syslog sourcetype = syslog
>	03/08/2020 12:04:38.000	Aug 3 12:04:38 ip-172-31-18-143 systemd[1]: Stopped target Cloud-config availability. host = ip-172-31-18-143 source = /var/log/syslog sourcetype = syslog

Result display type

Search Results

Field Sidebar

< Hide Fields

☰ All Fields

SELECTED FIELDS

a host 2
a source 3
a sourcetype 5

Default Selected Fields

a hold string value

INTERESTING FIELDS

a Ack 100+
a attack 100+
a class 22
Code 5
date_hour 14
date_mday 6
date_minute 60
a date_month 3
date_second 60
a date_wday 5
date_year 1
a date_zone 2
a dest_ip 100+
dest_port 100+

Interesting Fields

hold number value

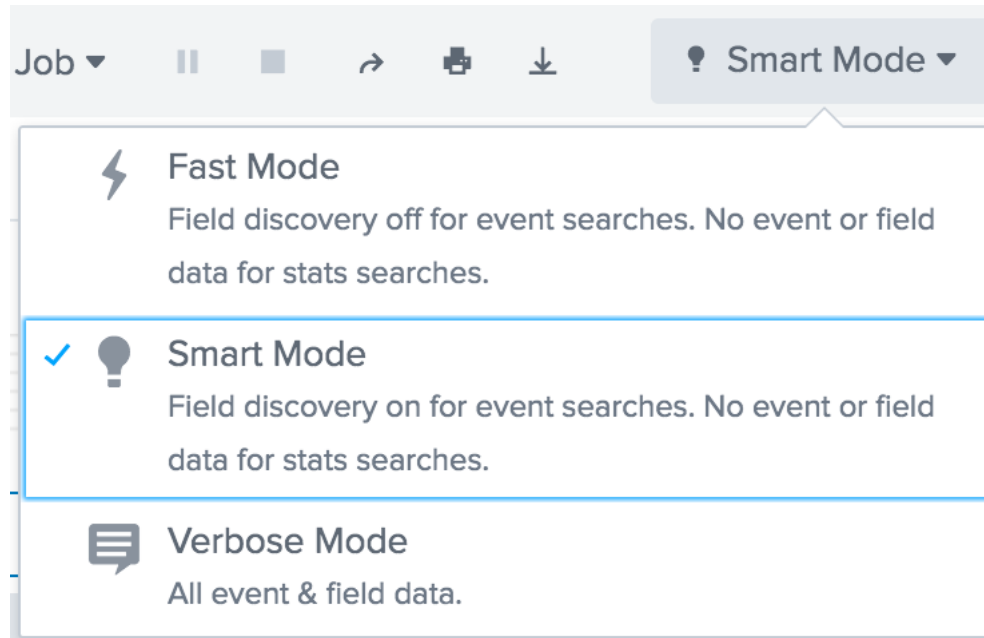
Selected Fields

host: A host is the name of the physical or virtual device where an event originates. It can be used to find all data originating from a specific device.

source: A source is the name of the file, directory, data stream, or other input from which a particular event originates.

sourcetype: Sources are classified into source types, which can be either well known formats or formats defined by the user.

Search Modes



Fast: Returns only Splunk Data like host, source, sourcetype, linecount, splunk_server, index.
and any specifically mentioned fields in search query

Verbose: Returns as much event information as possible, at the expense of slower search performance.

Smart: Toggles search behavior based on whether your search contains transforming commands or not.
For transforming searches, it behaves like Fast mode. For searches without transforming commands, it behaves like Verbose mode.


`_time` & `_raw`

The `_time` field contains an event's timestamp expressed in UNIX time (Epoch Time)

The `_raw` field contains the original raw data of an event.

1 index=se-assets

2 | table _time _raw

All time 

✓ 352,495 events (before 14/08/2020 17:57:56.000) No Event Sampling ▾

Job ▾ || ■ ➔ 🖨 ⬇ ⚡ Fast Mode ▾

Events Patterns **Statistics (352,495)** Visualization

100 Per Page ▾ ✎ Format Preview ▾

< Prev 1 2 3 4 5 6 7 8 ... Next >

<code>_time</code> ⬆	<code>_raw</code> ⬆	✎
2020-08-03 09:14:30	Aug 3 09:14:30 ip-172-31-18-143 sshd[8885]: Received disconnect from 222.186.15.115 port 63711:11: [preauth]	
2020-08-03 09:11:56	Aug 3 09:11:56 ip-172-31-18-143 sshd[8883]: Connection closed by 49.235.83.136 port 35968 [preauth]	
2020-08-03 09:05:18	Aug 3 09:05:18 ip-172-31-18-143 sshd[8878]: Disconnected from 222.186.175.23 port 10875 [preauth]	
2020-08-03 09:05:18	Aug 3 09:05:18 ip-172-31-18-143 sshd[8878]: Received disconnect from 222.186.175.23 port 10875:11: [preauth]	
2020-08-03 09:04:37	Aug 3 09:04:37 ip-172-31-18-143 sshd[8874]: Connection reset by 85.209.0.251 port 23522 [preauth]	
2020-08-03 09:04:35	Aug 3 09:04:35 ip-172-31-18-143 sshd[8876]: Connection reset by 85.209.0.251 port 23850 [preauth]	
2020-08-03 08:56:12	Aug 3 08:56:12 ip-172-31-18-143 sshd[8868]: Disconnected from 222.186.31.166 port 10630 [preauth]	
2020-08-03 08:56:12	Aug 3 08:56:12 ip-172-31-18-143 sshd[8868]: Received disconnect from 222.186.31.166 port 10630:11: [preauth]	

Basic Searches

SPL - Search Processing Language

Free-form search

Fields & Values - fields are case sensitive, values are not

Time Range Picker

Field Operators

Equal to

field = value

Not Equal to

field != value

Greater Than

field > value

Greater Than or Equal to

field >= value

Less Than

field < value

Less Than or Equal to

field <= value

Boolean Operators

AND – implied between terms, so you do not need to write it.

OR – used to specify that either one of two or more arguments should be true.

NOT – used to filter out events containing a specific word.

Wildcards

You can use the asterisk (*) character as a wildcard to match an unlimited number of characters in a string.

For example, "**fail***" matches **failure**, **failed**, **fails** and **failing**

clientip=10* matches all IPs starting with **10** like **100.** **101.** **102.** **10.**

to match 10.x.x.x network

clientip=10.*

Search Best Practices

Narrow the time window

Specify the index, source, or source type

Be specific

Avoid using NOT expressions

Filter as soon as possible

Apache Log Format

Log Format - %h %l %u %t \"%r\" %>s %b \"%{Referer}i\" \"%{User-agent}i\"

127.0.0.1 - frank [10/Oct/2000:13:55:36 -0700] "GET /apache_pb.gif HTTP/1.0" 200 2326 "http://www.example.com/start.html" "Mozilla/4.08 [en] (Win98; I ;Nav)"

Value	Std. Format	Description
127.0.0.1	%h	This is the IP address of the client (remote host) which made the request to the server
-	%l	The "hyphen" in the output indicates that the requested piece of information is not available. In this case, the information that is not available is the RFC 1413 identity of the client determined by identd on the clients machine.
frank	%u	This is the userid of the person requesting the document as determined by HTTP authentication.
[10/Oct/2000:13:55:36 -0700]	%t	The time that the server finished processing the request. The format is: [dd/mmm/yyyy:hh:mm:ss zone]
"GET /apache_pb.gif HTTP/1.0"	\"%r\"	The request line from the client is given in double quotes. The request line contains a great deal of useful information. First, the method used by the client is GET. Second, the client requested the resource /apache_pb.gif, and third, the client used the protocol HTTP/1.0
200	%>s	This is the status code that the server sends back to the client. This information is very valuable, because it reveals whether the request resulted in a successful response (codes beginning in 2), a redirection (codes beginning in 3), an error caused by the client (codes beginning in 4), or an error in the server (codes beginning in 5).
2326	%b	The last entry indicates the size of the object returned to the client, not including the response headers. If no content was returned to the client, this value will be "-".
"http://www.example.com/start.html"	\"%{Referer}i\"	This gives the site that the client reports having been referred from. (This should be the page that links to or includes /apache_pb.gif).
"Mozilla/4.08 [en] (Win98; I ;Nav)"	\"%{User-agent}i\"	The User-Agent HTTP request header. This is the identifying information that the client browser reports about itself.

Transforming Commands

search elements

| **command field1 field2**

search elements

| **command *function* field**

Transforming Commands

Command	Usage	Example
table	Gives the output in the form of a table	table field1 field2
rename	Rename a specific field	rename clientip as “source IP”
dedup	Removes duplicates	dedup clientip
sort	Sorts a field default ascending order	sort clientip sort – clientip sort clientip desc
top	Gives top fieldname default limit is 10	top clientip top limit=20 clientip
rare	Gives rare (least) occurring fieldname default limit is 10	rare clientip rare limit=5 clientip
head n	Returns latest ‘n’ results	head 5
tail n	Returns oldest ‘n’ results	tail 10

Stats Command

Function	Usage	Example
values	Lists the values of a specific field	stats values(status) by clientip
count	Gives the count of specific field	stats count BY clientip
dc distinct_count	Returns the number of unique values present in a field	stats dc(clientip)
sum	Sum of all the values of a field	stats sum(bytes)
max	Maximum individual value of a field	stats max(bytes)
min	Minimum individual value of a field	stats min(bytes)

Geo Location Commands

Function	Usage	Example
iplocation	Give Geo Location for all public IP s	iplocation clientip
geostat	Cluster map on map	geostat count by clientip